

Caden Cheng

408-592-6669 | cfc005@ucsd.edu | cadenfcheng.com | linkedin.com/in/cadenfcheng | github.com/cadencheng888

EDUCATION

University of California, San Diego

La Jolla, CA

B.S. in Computer Science and Mathematics, Data Science Minor

Expected Grad 2028

Relevant Coursework: Data Structures, Systems Programming, Algorithms, Operating Systems

EXPERIENCE

Software Engineer Intern

Jun. 2026 – Aug. 2026

Barobo, Inc

San Jose, CA

- Signed **CloudFront** video URLs with a configurable TTL (default **30 min**); rejects **~175** expired links/month
- Gated **10K** monthly video requests by account with per-user session and ownership records in **DynamoDB**
- Cut environment provisioning from 2h to 15m with **3 Terraform** modules and automated **AWS CLI** scripts
- Dropped p50 latency on the course-videos API **400→240ms** by collapsing **4 N+1 SQL** queries into joins

Software Developer → Incoming VP of Technology

Oct. 2025 – Present

Computer Science Engineering Society (CSES) at UCSD

La Jolla, CA

- Built **Google OAuth (NextAuth)** + **JWT** sessions for a lab-inventory app used by **100** members in **10** labs
- Enforced **4-role RBAC** across 12 API routes with 10 scoped permissions; 15 **Jest** tests run on every PR
- Rejected duplicate lab memberships with a unique compound index on (**user, lab**) in **Mongoose**

Software Developer

Mar. 2026 – Present

Triton Web Developers (TWD) at UCSD

La Jolla, CA

- Integrated **TinaCMS** so officers edit content without a PR; **40+** member edits shipped with **0** dev commits
- Modeled the homepage as **7** typed **CMS** sections, replacing hand-edited **HTML** that broke on every edit

PROJECTS

remark. [\[Link\]](#) | *Python, TypeScript, React, FastAPI, Claude API, Deepgram* | **Berkeley AI Hackathon** Jun. 2026

- Shipped in **24h** a voice agent turning overheard plans into calendar events and commands into actions
- Gave **Claude 5 tools** across **11** app categories, routed via **3 tiers**: Agentverse → Calendar → Browserbase
- Hit **3.2s** median speech-to-event latency (20 trials); **1.5s** of it is a deliberate silence-flush window
- Built an on-device **OpenCV** face gate (**36 ms/frame** Haar) so the mic streams only while a face is in view
- Streamed **Deepgram** transcripts over **WebSocket** with backoff reconnect; transcript buffer auto-wipes after **20s**

inVISION [\[Link\]](#) | *Python, OpenCV, MediaPipe, NumPy* | **Best Interactive AI, DiamondHacks** Apr. 2026

- Cut renderer peak memory **5x** (**34.7 → 6.7 MB** on a 267k-face mesh, tracemalloc) via **NumPy** transforms
- Mapped **21** MediaPipe hand landmarks to a smoothed rotation matrix: **23 ms/frame** detect, **~64 ms** max lag
- Profiled post-hackathon: per-face np.cross loop, not the transform, costs **664 ms/frame** on 10k faces
- Ran a **YOLOv8** detect → crop → background-removal pass on uploads before **TripoSr** mesh generation
- Built the **Flask** upload server and phone-camera web frontend so users capture the photo from their phone

Silent Speech [\[Link\]](#) | *PyTorch, MediaPipe, OpenCV* | **Best Hack, SanD Hacks**

Jan. 2026

- Trained a **BiGRU** on 90-frame clips of **88** mouth/jaw landmarks (of 478); **78%** (39/50) live word accuracy
- Noise + frame-drop augmentation raised val acc **72→79%** (3-seed avg); mixup dropped after no gain
- Recorded and labeled the 5-word landmark dataset in **24h**; normalized by mouth width for distance invariance
- Cut feature extraction **5x** (**0.25 → 0.05 ms/frame**) by dropping the 478-point mesh to the **88-point** subset

TECHNICAL SKILLS

Languages: Python, JavaScript, TypeScript, C/C++, C#, HTML/CSS, SQL

Frameworks & Domains: React, Next.js, Node.js, Tailwind CSS, PyTorch, TensorFlow, LLMs/GenAI, NLP, RAG

Developer Tools: Git/GitHub, Docker, CI/CD, AWS (CloudFront, DynamoDB, CLI), Terraform, Cursor, Gemini, Linux/Unix

Libraries: Pandas, NumPy, Matplotlib, Torchvision, OpenCV, MediaPipe, Mongoose